

Nifty Stock Price Prediction with LSTM

1. Introduction

The project aims to predict the Nifty50 stock index's closing prices using a Long Short-Term Memory (LSTM) neural network model. LSTM is a specialized type of Recurrent Neural Network (RNN) optimized for sequence prediction problems.

2. Data Acquisition and Features

2.1 Data Source

- Nifty50 Index: Historical stock market data representing the 50 most traded stocks of NSE India.

2.2 Data Features

- Date: Timestamp of recorded data.
- Close: Closing stock price for the day.

- **SMA_30: 30-day Simple Moving Average.**

- Formula: $SMA_t = \frac{1}{n} \sum_{i=t-n+1}^t P_i$

- **EMA_30: 30-day Exponential Moving Average.**

- Formula: $EMA_t = (Close_t - EMA_{t-1}) \times \frac{2}{(n+1)} + EMA_{t-1}$

3. Methodology

3.1 Data Collection

- Utilize the yfinance library to fetch historical data.
- Extract the "Close" prices to form the primary dataset.
- Calculate and append SMA_30 and EMA_30 to the dataset using their respective formulas.

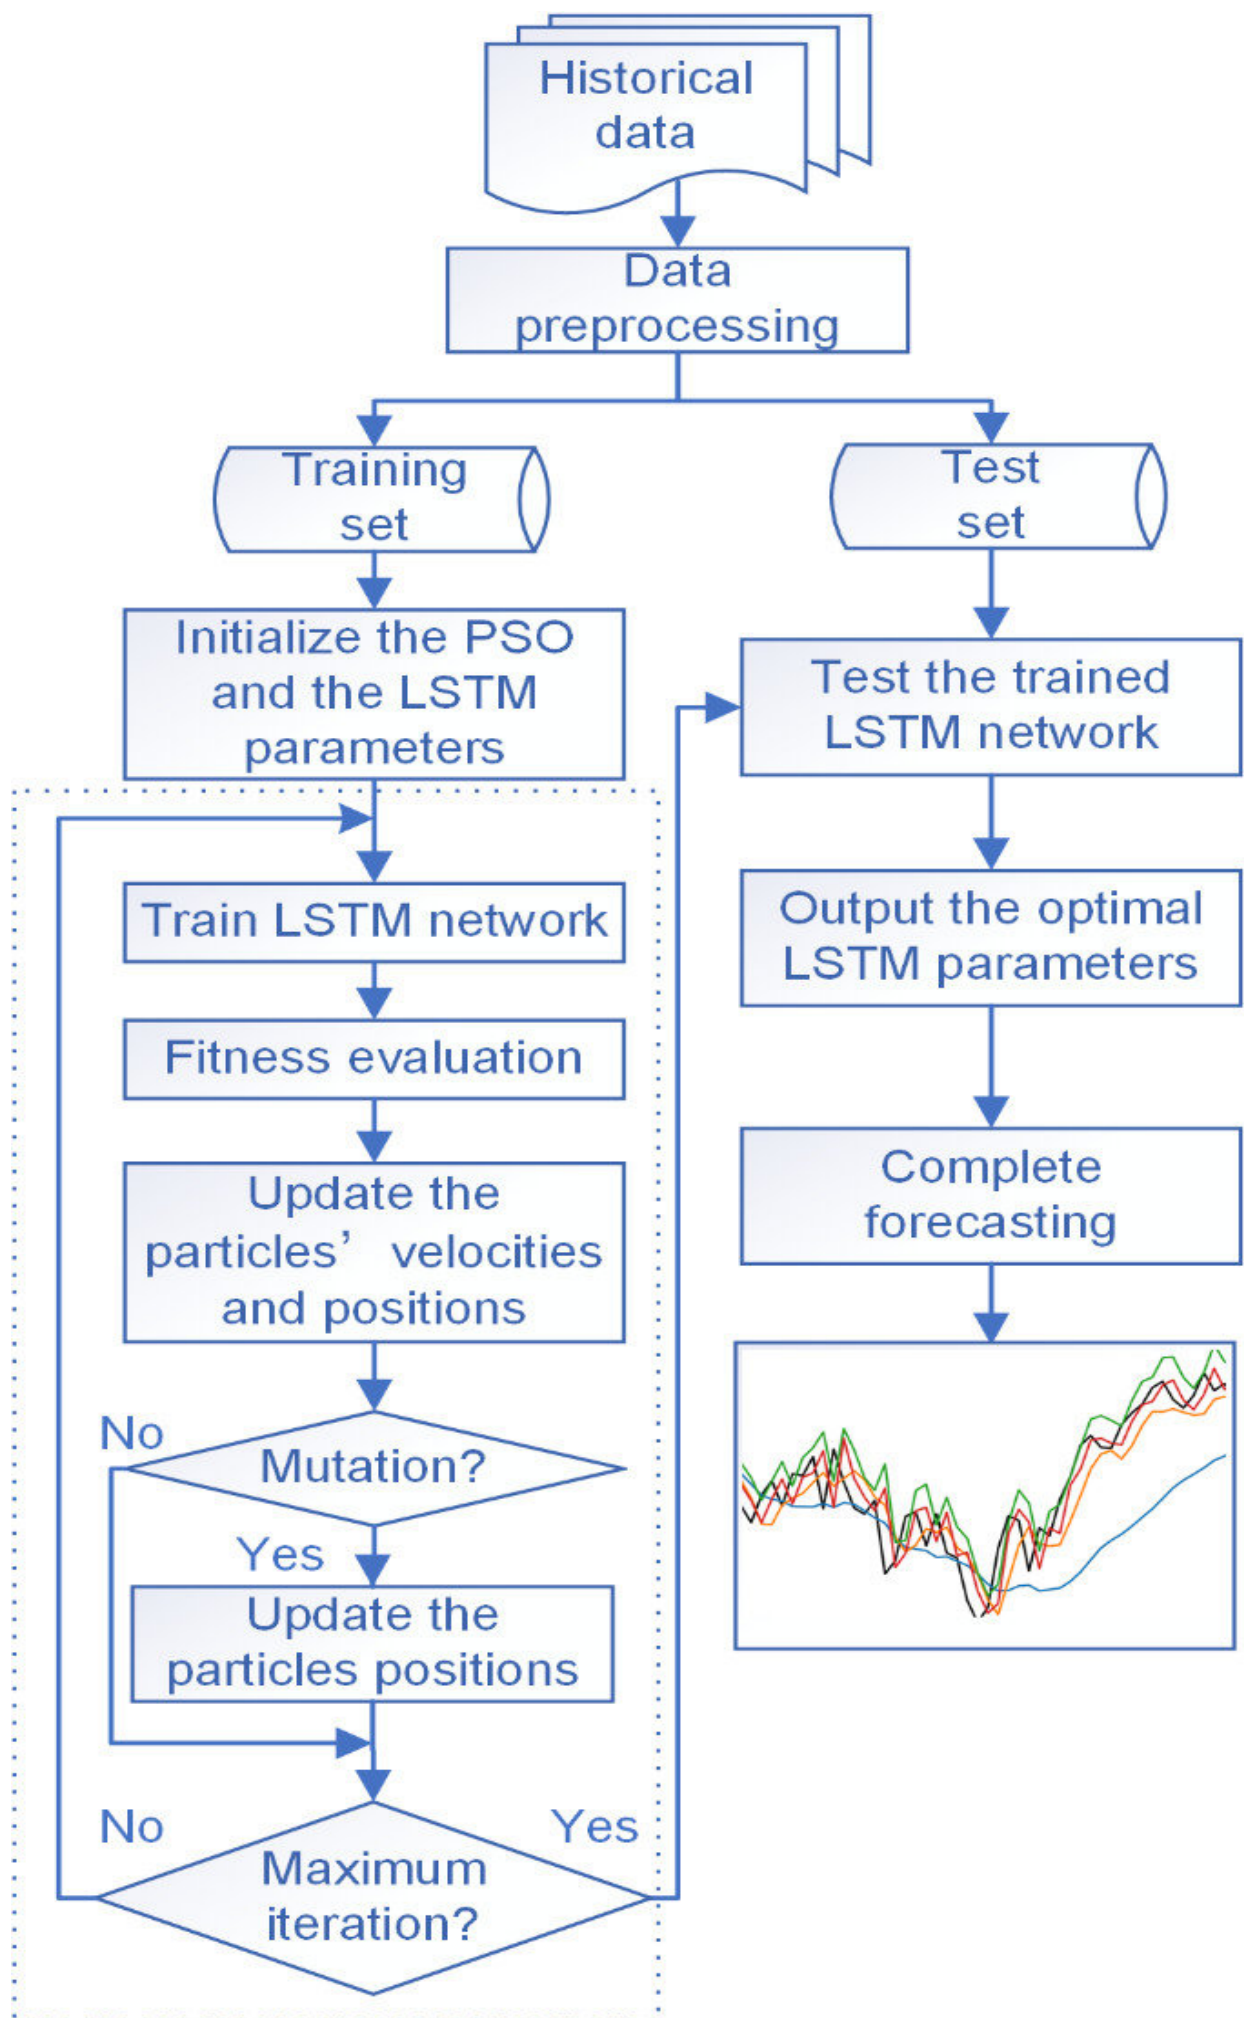
3.2 Data Preprocessing

3.2.1 Scaling

- Implement the MinMaxScaler from Scikit-learn.
- Transform the "Close" price data to fit within the range [0, 1].

3.2.2 Training and Test Data Preparation

- Allocate 75% of the scaled dataset for training the LSTM model.
- Utilize 25% for testing.
- Create overlapping sequences of 30 days to train the LSTM, as LSTM models require input sequences.



3.3 LSTM Model Architecture

3.3.1 Input Layer

- Input shape: (30, 1) representing a sequence of 30 days.

3.3.2 LSTM Layers

- First LSTM Layer:
 - 50 units.
 - Return sequences: True.
 - Activation function: tanh (hyperbolic tangent).
 - Recurrent activation: sigmoid.
- Second LSTM Layer:
 - 50 units.
 - Return sequences: False.
 - Activation function: tanh.
 - Recurrent activation: sigmoid.

3.3.3 Dropout Layers

- Dropout rate: 0.2 (20% of neurons are randomly turned off during each epoch).

3.3.4 Dense Layers

- First Dense Layer: 25 units with 'relu' activation.
- Output Dense Layer: 1 unit to predict the stock price.

3.4 Model Training

- Optimizer: Adam optimizer with adaptive learning rates.
- Loss Function: Mean Squared Error (MSE) to minimize the squared differences between actual and predicted values.
- Epochs: 200
- Batch Size: 50

3.5 Model Evaluation

- Implement Root Mean Square Error (RMSE) to gauge the model's error magnitude.
- R² (Coefficient of Determination) to measure how well future samples are predicted by the model.

3.6 Visualization

- Use the plotly library for interactive visualizations.
- Plot historical vs. predicted stock prices.
- Visualize the trend with SMA_30 & EMA_30.
- Forecast future stock prices.

4. Technical Implementation

4.1 Libraries and Dependencies

- yfinance: Fetch stock data.
- Keras: Neural network framework.
- Scikit-learn: Data preprocessing.
- Plotly: Data visualization.
- Pandas and NumPy: Data manipulation.

4.2 LSTM Mathematical Background

- LSTMs use three gates: input, forget, and output, each with its activation functions.
- The cell state in LSTM allows for long-term memory.

- Equations:

- Forget gate: $f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$
- Input gate: $i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$
- Update gate: $\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$
- Final cell state: $C_t = f_t \times C_{t-1} + i_t \times \tilde{C}_t$
- Output gate: $o_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$
- Hidden state: $h_t = o_t \times \tanh(C_t)$

